Local and Latent: Dual Views of Symmetry Breaking in Cooperative Multi-Robot RL

PI: Qi Zhang, Assistant Professor, Computer Science Department, Worcester Polytechnic Institute

Cash funding needed: $42,450; **AWS Promotional Credits needed:** $50,000

Amazon contacts: Zhenyu Zhang ⟨zhenyuzh@amazon.com⟩, Özalp Özer ⟨ozalpo@amazon.com⟩

Abstract Spatial symmetry is among the most effective inductive biases for robot learning, but in real multi-robot deployments it almost never holds exactly. Violations come in two structurally distinct forms: *local* breakers, visible in the current state but confined to small regions (e.g., fixed obstacles), and *latent* breakers, hidden from observation but fixed across an episode (e.g., wind or gravity). The two are orthogonal—state-dependent gating cannot find an unobserved breaker, and hidden-context inference cannot represent localized agent-dependent violations—yet realistic deployments exhibit both. We propose a dual-view theory and unified architecture. View 1 formalizes a locally factored partially symmetric Markov game and shows that critic-level symmetry error is controlled by discounted future breaker exposure, motivating a TD-learned per-agent gated critic. View 2 formalizes a latent symmetry-breaker POMDP and bounds value error by breaker-estimation error. The unified architecture pairs an equivariant latent encoder with per-agent gating.

Keywords cooperative multi-robot RL, group equivariance, symmetry breaking

1 Introduction

Group symmetries—geometric transformations leaving system properties invariant—provide one of the most effective inductive biases in robot learning, with consistent gains in imitation learning [1, 2, 3] and reinforcement learning (RL) [4, 5, 6], where exploration is expensive and demonstrations are scarce. Formally, given a group G acting on the state space $\mathcal{S}$ and action space $\mathcal{A}$ of an MDP, the MDP is G-symmetric if for every $g \in G$, $P(g \cdot s' | g \cdot s, g \cdot a) = P(s'|s, a), \quad R(g \cdot s, g \cdot a) = R(s, a)$. The notion extends to multi-robot settings by lifting actions and invariance to joint spaces (PI's [7]), where the relevant group factors as $G = G_{\text{spatial}} \times \mathcal{S}_n$. We focus on the spatial component and leave permutation symmetry for future work.

In practice, the symmetry condition almost never holds exactly. We observe that violations fall into two structurally distinct modes, separated by a single diagnostic: *is the breaker present in an agent's current observation, or hidden from it? Local* breakers are observable and confined to small regions of the per-agent state-action space—e.g., for warehouse robots near a fixed shelving unit, symmetry breaks only in the unit's vicinity, and at any moment only a subset of agents is affected. *Latent* breakers are hidden global variables fixed within an episode—e.g., wind on aerial drones, gravity on inclined multi-arm setups, current on underwater robots—where dynamics would be exactly symmetric *if* the vector were rotated together with state and action, yet the vector is not directly observable.

The two modes are structurally orthogonal: local methods cannot detect hidden breakers; latent inference cannot represent agent-dependent localized violations. Realistic deployments exhibit both. This *local-latent* dichotomy separates our proposal from prior work, which either treats single-agent settings where the dichotomy does not arise or globally relaxes symmetry without distinguishing structurally different modes [8, 9, 10, 11, 12, 13]. Closing this gap is the goal of this proposal.

2 Methods

We consider a discounted cooperative Markov game $\mathcal{M} = \langle \mathcal{N}, \mathcal{S}, \mathcal{A}, P, r, \gamma \rangle$ with agent set $\mathcal{N} = \{1, \ldots, N\}$, joint action space $\mathcal{A} = \mathcal{A}^1 \times \cdots \times \mathcal{A}^N$, and a spatial group G lifted to each agent's state

and action via group actions $(L_g^i, K_g^i)\colon \mathcal{X}^i \times \mathcal{A}^i \to \mathcal{X}^i \times \mathcal{A}^i$. We pursue two complementary views of why G may fail to be exact: View 1 attributes failures to localized regions of the observed local state-action space, motivating a state-dependent gated critic; View 2 attributes them to an unobserved global breaker variable (wind, gravity, current) and recovers equivariance in the augmented state through breaker inference. The two are combined in a single architecture.

2.1 View 1: symmetry breakers as locally observed state-action regions

We posit a *locally factored* partially symmetric Markov game in which the joint state, reward, and transition decompose agent-wise: $\mathcal{S} = \mathcal{X}^1 \times \cdots \times \mathcal{X}^N, r(s,a) = \sum_{i=1}^N r_i(x^i, a^i), P(\cdot \mid s,a) = \bigotimes_{i=1}^N P_i(\cdot \mid x^i, a^i)$. Each local factor carries its own finite symmetry group G_i acting on $(\mathcal{X}^i, \mathcal{A}^i)$ and its own *local breaker region* $\mathcal{B}^i \subseteq \mathcal{X}^i \times \mathcal{A}^i$. On the safe set $(x^i, a^i) \notin \mathcal{B}^i$ the local reward and transition are G_i-invariant; on $\mathcal{B}^i$ they may be arbitrary. The breaker region is itself G_i-invariant, so the safe/broken designation is a property of the orbit. This is more general than a single global gate over the joint state: symmetry can be exploited at agents whose local context is symmetric even when others' are not. We introduce a routed local model with an arbitrary routing parameter $\lambda_i(x^i, a^i) \in [0,1]$. Let $\overline{r}_i$ and $\overline{P}_i$ denote the G_i-orbit-averaged reward and transition. The routed factor interpolates between the equivariant and raw models: $\widetilde{r}_i^\lambda = (1-\lambda_i)\overline{r}_i + \lambda_i r_i$ and $\widetilde{P}_i^\lambda = (1-\lambda_i)\overline{P}_i + \lambda_i P_i$. Thus $\lambda_i = 0$ recovers the fully equivariant model, while $\lambda_i = 1$ recovers the raw model. The oracle current-breaker route $\lambda_i^{\mathrm{cur}}(x^i, a^i) = \mathbf{1}\{(x^i, a^i) \in \mathcal{B}^i\}$ is useful as a diagnostic, but our analysis and algorithm allow arbitrary learned routes.

Theory. For any bounded local critic Q_i, the one-step Bellman error of the routed model is controlled by the amount of equivariant averaging applied in nonsymmetric regions: $|(\widetilde{\mathcal{T}}_i^\lambda Q_i)(x,a) - (\mathcal{T}_i Q_i)(x,a)| \le (1-\lambda_i(x,a))(\Delta_i^r(x,a) + \gamma V_{i,\max}\Delta_i^P(x,a))$, where Δ_i^r and Δ_i^P are the reward and transition residuals between the raw and orbit-averaged local models. This bound captures the central tradeoff: small λ_i gives the sample efficiency of equivariance when residuals are small, while large λ_i avoids symmetry bias near local breakers. We will analyze how individual local critics' Bellman errors propagate to the joint critic and policy, showing that the local structure prevents error compounding across agents and that the policy gradient is similarly controlled by local Bellman errors. This justifies per-agent routing rather than a single global gate. The critic-level consequence is similar but more important for learning. Bellman contraction turns the one-step error into a corresponding fixed-point error, so a good route should suppress equivariant averaging precisely where it would introduce persistent value bias. Moreover, current breaker labels are not sufficient: even a currently safe pair $(x,a) \notin \mathcal{B}^i$ can have an asymmetric value if future trajectories enter a breaker region. We will formalize this through a discounted future-exposure term $\beta_i^{\pi_i}(x,a) = \mathbb{E}_{x,a}^{\pi_i}[\gamma^{\tau_i}]$, yielding the representative bound $|Q_i^{\pi_i}(L_g^i[x], K_g^i[a]) - Q_i^{\pi_i}(x,a)| \lesssim V_{i,\max}\beta_i^{\pi_i}(x,a)$. We will conduct a detailed analysis of the future-exposure bound, characterizing how it depends on the geometry of the breaker region and the policy's ability to avoid it, and showing that it can be much smaller than the current-breaker indicator.

Algorithm. The future-exposure bound motivates gating at the critic level by mixing two outputs: an orbit-averaged invariant critic $Q_\theta^{\mathrm{inv}}(s,a) = \frac{1}{|G|}\sum_{g\in G} f_\theta(g\cdot s, g\cdot a)$, an unconstrained raw critic $Q_\phi^{\mathrm{raw}}(s,a)$, and a learned gate $\widehat{\lambda}_\psi(s,a) = \sigma(g_\psi(s,a)) \in [0,1]$, combined as $Q_{\theta,\phi,\psi}^{\mathrm{gate}}(s,a) = (1-\widehat{\lambda}_\psi)\, Q_\theta^{\mathrm{inv}} + \widehat{\lambda}_\psi\, Q_\phi^{\mathrm{raw}}$. A local-factor variant sums per-agent gated branches, matching the future-exposure bound more directly. The gate is trained purely via TD loss—no auxiliary supervision on breaker labels. The gradient $\partial Q^{\mathrm{gate}}/\partial\widehat{\lambda}_\psi = Q_\phi^{\mathrm{raw}} - Q_\theta^{\mathrm{inv}}$ provides the routing signal: when the raw branch reduces TD error more than the invariant branch, the gate opens; when they agree, it favors the data-efficient invariant branch. We implement this with decentralized actors $a^i = \mu_i(o^i)$

and a centralized gated critic.

2.2 View 2: symmetry breakers as partially observed latent state

For hidden-context breakers, the agent-wise local-region view does not fit: the breaker is a single global object (e.g., a wind or gravity vector) shared by all agents and fixed across the episode. Throughout this subsection we therefore use single-agent-style notation (s, a, o, b), with the understanding that joint quantities $s = (x^1, \ldots, x^N)$, $a = (a^1, \ldots, a^N)$, $o = (o^1, \ldots, o^N)$ decompose per-agent and G acts simultaneously on all agents' frames; agent-wise factorization returns in §2.3. We model the environment as a *latent symmetry-breaker POMDP* $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{O}, \mathcal{B}, P_b, r_b, \Omega_b, \gamma, \mu_0, \rho_0 \rangle$ in which $b \sim \rho_0$ is sampled at the start of each episode and held fixed; the agent receives observations $o_t \sim \Omega_b(\cdot \mid s_t)$ and history h_t but never sees b. The structural assumption is *augmented equivariance*: for every $g \in G$, $r_{g\cdot b}(g{\cdot}s, g{\cdot}a) = r_b(s, a)$, $\quad P_{g\cdot b}(g{\cdot}C \mid g{\cdot}s, g{\cdot}a) = P_b(C \mid s, a)$, $\quad \Omega_{g\cdot b}(g{\cdot}D \mid g \cdot s) = \Omega_b(D \mid s)$. The system is symmetric if and only if b is transformed alongside (s, a); the visible state alone admits no group action consistent with the dynamics.

Theory. We will establish three results. First, if the dynamics are non-trivial in b, any critic enforcing observation-only equivariance $Q(g{\cdot}o, g{\cdot}a) = Q(o, a)$ is inconsistent with the true Bellman targets at some (o, a, g); the correct invariance is $Q(g{\cdot}o, g{\cdot}a, g{\cdot}z) = Q(o, a, z)$ where z is a breaker-sufficient statistic of history. Second, under Lipschitz assumptions $|r_b - r_{b'}| \leq L_r d_{\mathcal{B}}(b, b')$ and $\|P_b - P_{b'}\|_{\mathrm{TV}} \leq L_p d_{\mathcal{B}}(b, b')$, conditioning on an estimate $\widehat{b}$ rather than the true b yields a fixed-point value error bounded by $\frac{L_r + \gamma V_{\max} L_p}{1-\gamma}\, d_{\mathcal{B}}(b, \widehat{b})$, motivating auxiliary objectives that drive $d_{\mathcal{B}}(b, \widehat{b})$ down. Third, under (5) and an equivariant prior ρ_0, the Bayesian posterior over b transforms equivariantly: $\beta_t^{g\cdot h_t}(g \cdot B) = \beta_t^{h_t}(B)$; the encoder f_η should therefore be built so its output transforms in the same representation as b.

Algorithm. The architecture has three components. A latent breaker encoder $z_t = f_\eta(h_t)$ (recurrent or contrastive) is built so $f_\eta(g \cdot h_t) = g \cdot f_\eta(h_t)$, i.e., its output lives in the same representation as b. A breaker-conditioned policy and critic are equivariant in the augmented pair (o, z): $\pi_\theta(g{\cdot}a \mid g{\cdot}o, g{\cdot}z) = \pi_\theta(a \mid o, z)$ and $Q_\omega(g{\cdot}o, g{\cdot}a, g{\cdot}z) = Q_\omega(o, a, z)$. The encoder is trained with auxiliary objectives in addition to TD: a dynamics loss $\mathcal{L}_{\mathrm{dyn}} = \mathbb{E}[\ell(D_\psi(o_t, a_t, z_t), o_{t+1})]$, a contrastive loss pulling same-episode representations together (and pushing across-episode pairs apart), and an optional supervised loss $\mathcal{L}_{\mathrm{sup}} = \mathbb{E}[d_{\mathcal{B}}(f_\eta(h_t), b)^2]$ when simulator labels are available. We will explore three variants: a point-estimate encoder producing $\widehat{b}_t \in \mathcal{B}$; a belief representation outputting distributional parameters; and canonicalization, rotating inputs into a frame aligned with $\widehat{b}_t$, applying a standard network, and rotating outputs back.

2.3 Unifying perspective and combined architecture

The two views compose because they operate at different granularities: View 1 is per-agent and state-dependent; View 2 is global and per-episode. Concretely, we replace View 1's orbit-averaged invariant branch with View 2's breaker-conditioned variant—conditioning on a shared encoder $z_t = f_\eta(h_t)$ broadcast "tto all agents—while keeping the per-agent gate. Each agent's critic becomes $Q_i(x^i, a^i, h_t) = \big(1 - \widehat{\lambda}^i_\psi(x^i, a^i)\big)\, q^{\mathrm{inv}}_{i,\theta}(x^i, a^i, z_t) + \widehat{\lambda}^i_\psi(x^i, a^i)\, q^{\mathrm{raw}}_{i,\phi}(x^i, a^i, h_t)$, summed to $Q(s, a, h_t) = \sum_i Q_i$. Confirming that this combined architecture dominates both single-view variants on tasks mixing the two failure modes is a primary empirical claim.

Evaluation. We evaluate across three task families, each isolating a different combination of failure modes. *Multi-robot navigation with fixed obstacles* (cooperative goal-reaching on a 2D arena with a few fixed asymmetric obstacles) isolates agent-wise local breakers: at any timestep a subset of agents is near an obstacle while others are not; the group is $G = \mathrm{SO}(2)$ or a discrete subgroup, with symmetry-breaking strength varied by obstacle density. *Hidden wind navigation* replaces

the obstacles with a per-episode constant wind vector omitted from observations, isolating hidden global breakers; wind direction must be inferred from the discrepancy between commanded and realized displacement. Finally, *combined regime* is multi-arm manipulation on an inclined surface (hidden gravity) with localized obstacles, testing the combined architecture and the orthogonality claim.

Metrics follow standard practice—final return, sample efficiency, training stability across seeds—plus view-specific diagnostics: gate activation maps and their correlation with oracle breaker labels (View 1), and angular error between $\widehat{b}_t$ and the true wind/gravity vector (View 2). Performance is reported as a function of symmetry-breaking strength to characterize the regime where each view, and the combined architecture, dominates.

Baselines and ablations. On View 1 we compare against a raw MLP critic, a fully invariant critic, and a *current-breaker oracle gate* using ground-truth $\mathbf{1}\{(s,a) \in \mathcal{B}\}$—a diagnostic for whether current breaker information alone suffices (the future-exposure bound predicts it does not when trajectories cross breaker regions several steps ahead). On View 2: an unstructured MLP, a strict observation-only equivariant model, a recurrent non-equivariant model with history but no symmetry prior, an *oracle* receiving the true b (upper bound), and an *observable-breaker ablation* exposing b in the observation (confirming the latent variable is the missing information). We ultimately will compare the combined architecture against the strongest single-view variant from each.

3 Expected results: key outcomes

The project spans 14 months. The first three months focus on setup and baselines implementation. View 1 theory and algorithms are developed in months 2–6, running in parallel with View 2 theory and algorithms in months 4–8. Unification of the two views begins in month 8 and continues through month 11. The final three months are dedicated to writing code, paper, and submission.

Funds needed. *Cash funding:* \$42,450 is requested for the annual stipend of a research assistant (RA) [14]. *AWS promotional credits:* \$50,000 is requested. ∼\$35,000 will be allocated to EC2 for parallelized multi-robot RL training, ∼\$10,000 to RoboMaker and SageMaker to enable realistic and large-scale robotics simulation, training, and benchmarking, ∼\$5,000 to supporting infrastructure.

References

[1] Andy Zeng, Pete Florence, Jonathan Tompson, Stefan Welker, Jonathan Chien, Maria Attarian, Travis Armstrong, Ivan Krasin, Dan Duong, Vikas Sindhwani, et al. Transporter networks: Rearranging the visual world for robotic manipulation. In *Conference on Robot Learning (CoRL)*, 2021.

[2] Dian Wang, Stephen Hart, David Surovik, Tarik Kelestemur, Haojie Huang, Haibo Zhao, Mark Yeatman, Jiuguang Wang, Robin Walters, and Robert Platt. Equivariant diffusion policy. In *Conference on Robot Learning (CoRL)*, 2024.

[3] Joohwan Seo, Soochul Yoo, Junwoo Chang, Hyunseok An, Hyunwoo Ryu, Soomi Lee, Arvind Kruthiventy, Jongeun Choi, and Roberto Horowitz. Se(3)-equivariant robot learning and control: A tutorial survey. *International Journal of Control, Automation and Systems*, 23.

[4] Elise Van der Pol, Daniel Worrall, Herke van Hoof, Frans Oliehoek, and Max Welling. Mdp homomorphic networks: Group symmetries in reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pages 4199–4210, 2020.

[5] Dian Wang, Robin Walters, and Robert Platt. So(2)-equivariant reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2022.

[6] Jinzhu Luo, Dingyang Chen, and Qi Zhang. Reinforcement learning with euclidean data augmentation for state-based continuous control. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL https://openreview.net/forum?id=NwiFLtWGEg.

[7] Dingyang Chen and Qi Zhang. E (3)-equivariant actor-critic methods for cooperative multi-agent reinforcement learning. In *Proceedings of the 41st International Conference on Machine Learning*, pages 7640–7666, 2024.

[8] Marc Finzi, Gregory Benton, and Andrew Gordon Wilson. Residual pathway priors for soft equivariance constraints. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pages 30037–30049, 2021.

[9] Rui Wang, Robin Walters, and Rose Yu. Approximately equivariant networks for imperfectly symmetric dynamics. In *International Conference on Machine Learning (ICML)*, pages 23078–23091, 2022.

[10] David W Romero and Suhas Lohit. Learning partial equivariances from data. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 36466–36478, 2022.

[11] Tycho van der Ouderaa, David W Romero, and Mark van der Wilk. Relaxing equivariance constraints with non-stationary continuous filters. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 33818–33830, 2022.

[12] Elyssa Hofgard, Rui Wang, Robin Walters, and Tess Smidt. Relaxed equivariant graph neural networks. In *ICML Workshop on Geometry-grounded Representation Learning and Generative Modeling (GRaM)*, 2024. Extended Abstract.

[13] Jung Yeon Park, Sujay Bhatt, Sihan Zeng, Lawson LS Wong, Alec Koppel, Sumitra Ganesh, and Robin Walters. Approximate equivariance in reinforcement learning. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2025.

[14] Worcester Polytechnic Institute. Stipend and hourly rates, 2026. URL https://www.wpi.edu/offices/talent/student-employment/graduate-worker-resources/stipend-and-hourly-rate. Accessed: 2026-05-11.